SEED ROUND · 2026

# Mesh.

The AI-native coding environment — web app & desktop.
Voice-driven. Compression-first.

**Edgar Baumann**
Co-Founder · WU Wien

**Philipp Horn**
Co-Founder · WHU

| | | |
|---|---|---|
| **74%** | **3.9×** | **€500k** |
| avg token reduction (Capsule) | more codebase per context window | seed ask · 18mo runway |

edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com

THE PROBLEM

# AI coding tools are hitting a wall — and it's called context.

Every AI tool faces the same constraint: LLM context windows are finite, and real codebases are not.

- **AI assistants hallucinate** because they only see a fraction of the codebase at once
- **API costs explode** as teams load more code — paying for tokens, not for answers
- **Large files are the worst offenders** — a single 50 KB module consumes most of a 128k window
- **Voice is completely untapped** — developers still type every command, even when speaking is 3× faster

> *Sending raw code to an LLM is like faxing a dictionary to answer a single question. You're paying for 99% noise to get 1% signal.*

The average developer spends 42% of their time understanding code, not writing it. *(Stack Overflow Developer Survey 2023)*

THE SOLUTION

# One environment. Full context. Voice + AI.

Mesh combines three capabilities no other tool has together — as a web app and desktop app.

**74%**

**Capsule Compression**

Proprietary structural compression — 3.9× more codebase per context window, 75% lower API costs, better accuracy

⌨

**Unified Workbench**

Editor + Terminal + AI chat + Dependency Graph — one surface, zero context switching

🎙

**Voice-Driven Agent**

Speak your intent. Mesh reads compressed codebase context, generates code changes, explains them

> *The only tool that solves **context, cost, and workflow** at the same time.*

THE TECHNOLOGY

# The engine that makes everything else possible.

Capsule compresses source code structurally before it reaches any LLM. Not summarization — the full file is always recoverable on demand.

## Why context quality matters

### A — NIAH (Needle In A Haystack)

Models degrade at >60–70% context utilization — "Lost in the Middle" effect (Liu et al., 2023). At 100k+ tokens, retrieval accuracy drops 20–40% based on position in context. **Capsule keeps utilization low regardless of codebase size.**

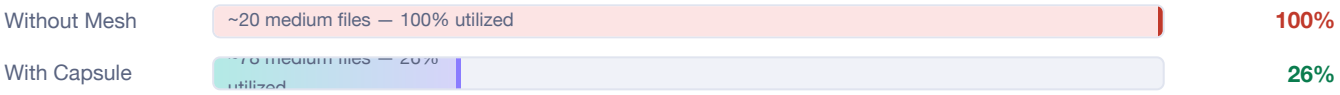


### B — SWE-bench

Top agents (Claude Opus 4.6, GPT-4o, Gemini) solve ~40–55% of real GitHub issues. The leading failure mode: **insufficient codebase context**. More files in context = fewer failures from missing information.

## C — Capsule benchmark results

| FILE SIZE | TOKEN REDUCTION | SAVING |
|---|---|---|
| ~200B (xs) — pass-through | | +10% |
| ~1 KB (small) | | –83% |
| ~5 KB (medium) | | –95% |
| ~18 KB (large) | | –98.5% |
| ~50 KB (XL) | | –99.4% |
| ~100 KB (XXL) | | –99.9% |
| **Average (small–XXL)** | | –74% → 3.9× |

PRODUCT

# Built for the way developers actually think.

Three surfaces, one flow. Web app + desktop — no install required for the web version.

| SURFACE | WHAT IT DOES |
|---|---|
| **Editor** | Monaco-based editor with AI chat panel, file explorer, live dependency graph, and workspace intelligence sidebar. The AI knows your entire codebase before you ask — compressed, indexed, always current. |
| **Terminal** | Dedicated terminal workspace — not a bottom panel but a full surface. Runs alongside the editor without losing context. |
| **Voice-Coding** | Speech-driven agent. Say "*refactor the auth middleware to use JWT*" — Mesh reads the relevant files via Capsule, generates the change, and explains it. No typing required. |

## Key differentiators

- **Capsule compression** — 74% fewer tokens, same intelligence, 75% lower API cost
- **Live dependency graph** — updates in real time as code changes
- **Persistent workspace memory** across sessions
- **Multi-model** — Claude Opus/Sonnet, Gemini, GPT — no provider lock-in

MARKET OPPORTUNITY

# Every developer is a potential user.

We're starting with those who feel the API cost pain most — indie developers and small teams.

**$28B**
TAM — Global developer tools (2024, 12% CAGR)

**$4.2B**
SAM — AI-enhanced IDE / coding assistant segment (2025)

**$120M**
SOM (3yr) — Indie devs + small teams, Europe + North America

## Why now

- LLMs crossed the threshold where voice-to-code is genuinely usable
- Cursor proved developers pay for AI-native editors — **~$400M ARR in 2 years at $20/month**
- API costs are a real budget line item for dev teams — compression has measurable ROI
- Green tech pressure: enterprises increasingly track and report AI energy spend


edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com

COMPETITIVE LANDSCAPE

# We're not an extension. We're the environment.

| | MESH | VS CODE + COPILOT | CURSOR | GOOGLE ANTIGRAVITY |
|---|---|---|---|---|
| Voice-driven coding agent | ✓ | ✗ | STT only [2] | ✗ |
| Structural token compression | ✓ | ✗ | ✗ | ✗ |
| Whole-codebase AI context | ✓ | Partial [1] | Partial [1] | Partial [1] |
| Web app (no install) | ✓ | Partial [3] | ✗ | ✗ |
| AI-native architecture | ✓ | ✗ (retrofitted) | ✓ | ✓ |
| Multi-model | ✓ | ✓ | ✓ | ✓ |

[1] Embedding retrieval — finds relevant files, sends them raw to the model

[2] Cursor has speech-to-text input (Ctrl+M dictation into chat) — not a voice-driven coding agent

[3] VS Code for the Web exists but has no terminal, no debugger, and limited extension support

> *All three retrieve code with embeddings and send it raw — cost scales with codebase size.* ***Capsule compresses before the model sees it — cost stays flat.***

BUSINESS MODEL

# Freemium → Pro → Teams

**FREE**

**€0**
forever

- Full editor + terminal
- AI chat (limited requests/mo)
- Basic workspace indexing

**PRO — MOST POPULAR**

**€19**
per month

- Unlimited AI requests
- Full Capsule compression
- Voice-coding agent
- Priority model access

**TEAMS**

**€49**
per seat / month

- Everything in Pro
- Shared workspace intelligence
- Team codebase context
- Admin controls + analytics

> *Cursor charges $20/month and crossed $400M ARR. Mesh saves time and API costs — a double value proposition that compounds at the team tier.*


edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com

TRACTION

# Early stage — real foundation.

- **MVP complete** — Editor, Terminal, and Voice-Coding surfaces are fully functional
- **Capsule pipeline live and benchmarked** — 74% average token reduction across all file types and sizes
- **Architecture built for scale** — gateway/worker split, multi-provider AI support, horizontal scaling
- **Active development velocity** — 8 shipped phases with verified completion

NEXT MILESTONE

### First 100 beta users

Measure retention, voice feature engagement, and real-world token savings to validate the core thesis


edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com

TEAM

# Built by developers, for developers.

### Edgar Baumann

**Co-Founder · WU Wien**

Built Mesh end-to-end — architecture, compression pipeline, voice agent, and frontend. Driven by the gap between how developers think and how their tools make them work.

### Philipp Horn

**Co-Founder · WHU**

Strong entrepreneurial instinct — business strategy, GTM, and growth.

> *Mesh comes from felt pain, not a market analysis. We built the tool we needed, and it didn't exist.*


edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com

VISION

# In 5 years, sending raw code to an AI will seem as wasteful as printing emails.

| HORIZON | GOAL |
|---|---|
| **12 months** | 1,000 paying Pro users · Capsule as the industry reference benchmark · first team accounts |
| **24 months** | 10,000 users · Series A · open Capsule SDK for third-party integrations |
| **5 years** | The default environment for developers who think faster than they type — plus a compression infrastructure layer powering other AI tools |

SEED ROUND

# €500,000 · 18-month runway

- Engineering — 2 senior engineers, voice pipeline + Capsule v2
- Growth — developer marketing, OSS, content
- Infrastructure — AI inference at scale
- Operations — legal, tools

edgar.baumann@try-mesh.com · philipp.horn@try-mesh.com · Demo available on request